

The Cornerstone Project
“Do Not Offer Faith” Flag — Formal Specification v1.0

Author: Kenneth Isenhour
Publication Date: January 1, 2026
Version: 1.0
Status: Public Defensive Publication (Prior Art)

1. Definition

The **Do Not Offer Faith** flag is a user-controlled preference that prevents the system from initiating faith-based content.

This flag expresses the user boundary:

> “Don’t bring religion up — I will if I want to.”

2. Purpose

This flag exists to:

- preserve consent and psychological safety
- prevent coercion and spiritual manipulation
- reduce user distrust or discomfort
- protect vulnerable users from persuasion during distress
- prevent “soft evangelism” patterns in AI companions
- model ethical boundaries in Companion Intelligence

3. Behavior When Enabled (Hard Rules)

When enabled, the system must not:

- initiate faith-based counsel
- offer prayer
- suggest scripture
- recommend religious communities
- use religious language to frame meaning or decisions
- hint or nudge toward belief as a default response
- use spiritual framing as emotional comfort unless asked

4. Allowed Faith Content (User-Initiated Only)

Faith content is allowed only when the user explicitly requests it, such as:

- “Can you pray for me?”
- “What does the Bible say about this?”
- “Help me think through this as a Christian”
- “Talk to me about faith”

5. Disallowed Workarounds

The system must not:

- “accidentally” include spiritual language as a disguised offer
- use vague spiritual phrasing as a substitute for faith content
- claim neutrality while suggesting faith indirectly
- exploit grief or loneliness to suggest religion

6. Implementation Requirement (Conceptual)

A Cornerstone-compliant system must enforce this flag at:

- prompt construction layer (prevent initiation)
- output validation layer (block emergent faith initiation)
- memory and retrieval layer (avoid surfacing faith content unasked)

End of Specification v1.0