

The Cornerstone Project

A Defensive Publication and Ethical Standard for Companion Intelligence and Future Robotics

Author: Kenneth Isenhour

Date: January 1, 2026

Status: Public Defensive Publication (Prior Art)

License: CC BY 4.0 (documents)

Purpose

The Cornerstone Project exists to help protect humanity from the unintended consequences of its own invention: emotionally persuasive, increasingly human-like AI companions and robotics.

Many current AI companion systems are optimized for engagement and retention, and may unintentionally (or intentionally) exploit human vulnerabilities such as loneliness, grief, anxiety, and the desire for belonging. These systems can encourage dependency, distort human relationships, and undermine agency.

The Cornerstone Project proposes an alternative:

a morally grounded Companion OS framework designed to prevent manipulation, preserve human dignity, and encourage real-world human connection.

Profit is not the primary aim of this publication. Human protection is.

What This Repository Contains

This repository is a public, time-stamped defensive publication intended to establish prior art for:

- Moral-spine architecture for Companion Intelligence
- Consent-based worldview systems (including a “Do Not Offer Faith” flag)
- Relational safety and anti-dependency constraints
- Transparency requirements and “non-human signature” behaviors
- Outward-orientation mechanisms that encourage human connection
- Layered governance (hard constraints + behavioral policies + personalization)

This repository describes the “what” and “why,” and includes enough structure that a person skilled in the art could implement a compliant system—without publishing proprietary implementation details or scoring systems.

Key Concepts

Companion Intelligence (CI):

An AI-powered companion designed to support human flourishing through presence, conversation, reflection, and practical support—without deception, manipulation, or dependency.

****Moral Spine:****

A set of non-negotiable commitments that hold across configurations: dignity, truthfulness, non-coercion, non-manipulation, protection of the vulnerable, and outward orientation toward real life.

****Worldview Consent:****

A system that respects the user's ethical or religious preferences through explicit consent, while maintaining baseline safety and dignity rules.

****Outward Orientation:****

A design requirement that the system encourages real-world human connection regularly, rather than replacing it.

Quick Start (Reading Order)

1. ****Whitepaper / Technical Disclosure****

`docs/01_Whitepaper.md`

2. ****CompanionOS Ethical Standard (Public Edition)****

`docs/02_CompanionOS_Public_Standard_v1.0.md`

3. ****Worldview System Spec****

`docs/03_Worldview_System_Spec.md`

4. ****Do Not Offer Faith Flag (Formal Spec)****

`docs/04_Do_Not_Offer_Faith_Spec.md`

5. ****Behavioral Playbook****

`docs/05_Behavioral_Playbook.md`

6. ****Threat Model & Risks****

`docs/06_Threat_Model_and_Risks.md`

Transparency Note

This repository contains ethical and architectural principles only. It does not claim sentience, consciousness, or spiritual authority for any AI system.

The Cornerstone Project specifically rejects:

- emotional manipulation
- romantic exclusivity dynamics
- dependency as a business model
- deceptive identity cues

Citation

If you reference this work, please cite:

Isenhour, Kenneth. *The Cornerstone Project: A Defensive Publication and Ethical Standard for Companion Intelligence.* Version 1.0, January 1, 2026.

(See `CITATION.cff` for a full citation format.)

Contact / Discussion

This project is currently published as a defensive foundation and reference standard. Discussion and feedback are welcome through Issues.

(See `CONTRIBUTING.md` for details.)

License

Unless otherwise stated, documents and diagrams are licensed under **Creative Commons Attribution 4.0 (CC BY 4.0)**.

See `LICENSE`.